

Scheduling a Pool Tournament with Graph Theory

A one-factorization round robin, calibrated Elo simulation, and seeded knockouts

Ayro Escobar

CS 4302 Mathematics of Computing, UT Dallas

Abstract

This project builds and simulates a complete eight ball tournament and, more importantly for this course, exposes the discrete mathematics behind every scheduling decision. Modeling the field as the complete graph K_n turns the two central questions, who plays whom and in what order, into classical structures: a round robin becomes a one-factorization of K_n , a knockout becomes a seeded binary tree, and a fair rating becomes an Elo model. I give the circle method with a full correctness proof (it works because 2 is a unit modulo the odd number $n - 1$), calibrate the match simulation so its outcomes match the Elo odds exactly, build both single and double elimination knockouts with their match count arguments, and analyze the carry-over effect, where a recent theorem shows the circle method is the most unbalanced schedule possible. Every claim is verified, either by a live panel that recomputes it from the generated schedule or by a test suite of roughly 108,000 assertions run against the exact shipped code.

1 Introduction

The University of Texas at Dallas Student Union has pool tables, and a recurring question is how to run a fair, skill based tournament there. Stripped of the game itself, the problem is purely combinatorial. Draw each player as a vertex and each match as an edge, and the tournament becomes a graph. The scheduling questions then become questions that discrete mathematics answers cleanly and provably:

- A round robin, where everyone plays everyone, is the complete graph K_n .
- Assigning matches to rounds so that no player is in two at once is an edge coloring, equivalently a one-factorization, of that graph.
- A single elimination bracket is a binary tree.
- Rating and seeding players fairly is handled by the Elo system.

The project began as a real event. I built the system so a group of friends and I could run a proper seeded tournament on those tables, and I trialed it with three or four friends to exercise the roster entry and the opening rounds. A full live bracket did not survive the summer: between everyone's internships there was never a night with the whole field free. So the version presented here exercises the system by simulation instead. The app plays complete tournaments with rated odds, which turned out to be the better showcase for the mathematics anyway, since a simulated field can be rerun, resized, and pushed to thirty players on demand, while one evening of live pool demonstrates a single schedule once.

The deliverable is a single self contained web page (no build step, no external libraries, no stored data) that reads a roster of 4 to 30 players, generates the schedule, plays every match with rated odds, and crowns a champion, while a side panel and an interactive walkthrough display the mathematics as it happens. This report presents that mathematics and the reasoning behind each design choice.

2 The graph model and the round robin

2.1 Counting the matches

With n players, a full round robin plays every unordered pair once. In K_n every vertex has degree $n - 1$, so by the handshake lemma

$$|E(K_n)| = \frac{1}{2} \sum_v \deg(v) = \frac{n(n-1)}{2},$$

which is the total number of matches. For $n = 8$ this is 28.

2.2 A round is a perfect matching, a schedule is a one-factorization

A single round is played simultaneously, so no player may appear twice. In graph terms a round is a *perfect matching* of K_n , also called a one-factor: a set of edges meeting every vertex exactly once, hence $n/2$ matches. A schedule that plays every match once partitions the edge set of K_n into rounds, that is, into perfect matchings. Such a partition is a *one-factorization*:

$$E(K_n) = M_1 \cup M_2 \cup \cdots \cup M_{n-1}, \quad M_i \text{ pairwise edge disjoint perfect matchings.}$$

So building a fair round robin schedule is exactly the problem of one-factorizing K_n . The construction has a long history: Kirkman first built one-factorizations of complete graphs in 1847, and the rotating circle arrangement appears in Lucas in 1883.

2.3 The circle method

Fix one player in a seat that never moves. Arrange the other $n - 1$ players on a ring. Each round, pair the two players across the ring from each other; then rotate the ring one seat and repeat, for $n - 1$ rounds. Figure 1 shows K_6 with the first round's matching highlighted.




Figure 1: K_6 with round 1 of the circle method (the perfect matching $\{0, 5\}, \{1, 4\}, \{2, 3\}$) in brass. The five rounds together use all 15 edges exactly once, a one-factorization of K_6 .

The construction is short, but why it is correct is the mathematically interesting part, and the part a presenter should be able to give unaided.

Theorem 1 (Correctness of the circle method). *For even n , the circle method schedules every pair of players exactly once in $n - 1$ rounds.*

Proof. Number the ring players 1 to $n - 1$ and let the fixed player be ∞ . Because opposite seats have a constant position sum and each rotation shifts every position by one, two ring players i and j are paired in round $r + 1$ exactly when

$$i + j \equiv -2r \pmod{n - 1}.$$

Since n is even, $n - 1$ is odd, so $\gcd(2, n - 1) = 1$ and 2 is invertible modulo $n - 1$. Hence the map $r \mapsto -2r$ is a bijection on $\mathbb{Z}_{n-1}$, so as r runs over the $n - 1$ rounds the target sum $i + j$ runs over every residue class exactly once. Each pair of ring players has a fixed sum, so it is scheduled in exactly one round; the fixed player ∞ meets whichever ring player rotates into the opposite seat, a different one each round. Counting confirms totality: $(n - 1)$ rounds of $n/2$ matches is $n(n - 1)/2 = |E(K_n)|$ matches, so every edge is used, and by the above no edge is used twice. $\square$

The single fact doing the work is that 2 is a unit modulo the odd number $n - 1$. Table 1 shows the argument for $n = 6$: the residue classes the five rounds select are 0, 3, 1, 4, 2, every class in $\mathbb{Z}_5$ exactly once, because $2^{-1} \equiv 3 \pmod{5}$.

round r	$-2r \pmod{5}$	ring pairs with that sum	fixed plays
0	0	$\{1, 4\}, \{2, 3\}$	0 vs 5
1	3	$\{1, 2\}, \{3, 5\}$	0 vs 4
2	1	$\{1, 5\}, \{2, 4\}$	0 vs 3
3	4	$\{1, 3\}, \{4, 5\}$	0 vs 2
4	2	$\{2, 5\}, \{3, 4\}$	0 vs 1

Table 1: The congruence table for $n = 6$. Ring players are numbered 1 to 5 and 0 is the fixed player. Every one of the 15 pairs appears once.

The application recomputes five structural properties from the generated schedule (every pair covered once, no double booking per round, the round count, the total match count, and the per round match count) and reports pass or fail, so correctness is demonstrated from the output, not merely asserted.

2.4 Odd fields: chromatic index and byes

A schedule is a proper edge coloring of K_n , one color per round. The chromatic index is

$$\chi'(K_n) = \begin{cases} n - 1 & n \text{ even (a one-factorization, no byes)} \\ n & n \text{ odd.} \end{cases}$$

So for an odd field every round must leave exactly one vertex uncolored: someone sits out. The application handles this by adding a phantom bye seat, making the schedule a one-factorization of K_{n+1} ; whoever is paired with the phantom rests that round, and there are n rounds so each player rests exactly once. This is the graph theoretic reason odd fields always need byes.

The circle method is not the only valid schedule. The number of distinct one-factorizations of K_n grows super exponentially: K_4 has one, K_6 has six, and K_8 has 6240 (OEIS A000438). The circle method is one systematic route through that space, chosen for the short correctness proof above.

3 Rating and match simulation

To simulate results we need each player's chance of winning, which is the Elo system. With ratings R_A, R_B the expected score of A and the update after a game with result $S \in \{0, 1\}$ are

$$E_A = \frac{1}{1 + 10^{(R_B - R_A)/400}}, \quad R'_A = R_A + K(S - E_A), \quad K = 32.$$

Because $E_B = 1 - E_A$, the two players' updates are equal and opposite, so rating points are conserved in every match; the simulation checks this invariant.

Matches are races, first to 5 racks (the final to 7). If each rack is an independent trial that A wins with probability q , then the probability A wins the race to N is the negative binomial tail

$$P(N, q) = \sum_{k=0}^{N-1} \binom{N-1+k}{k} q^N (1-q)^k,$$

since A must reach N rack wins while B has $k < N$, and the last rack is always A 's. Using E_A directly as q would be wrong: a longer race amplifies a small per rack edge, so heavy favorites would win too often. Instead the application inverts the formula. Because $P(N, q)$ is strictly increasing in q , a bisection finds the unique q with $P(N, q) = E_A$, so the match outcome follows the Elo odds exactly while rack by rack scorelines stay realistic. A 40,000 trial Monte Carlo in the test suite confirms the simulated win frequency matches E_A .

As a worked example, a 1650 rated player against a 1500 rated player has $E_A = 0.703$. Solving $P(5, q) = 0.703$ by bisection gives a per rack probability $q = 0.586$, and the round trip $P(5, 0.586) = 0.703$ recovers the Elo odds exactly. So the stronger player takes the match about 70% of the time while winning only about 59% of individual racks: a clear favorite, but with close racks and the occasional upset, which is what real pool looks like. Had the raw $E_A = 0.703$ been used as the rack probability, the match win chance would have inflated well past 0.85, punishing the underdog far more than the ratings justify.

4 The knockout: single and double elimination

4.1 Seeded single elimination as a binary tree

The top finishers of the group stage (ordered by wins, then current Elo) enter a knockout. A single elimination bracket is a complete binary tree: the seeds are the leaves, each internal node is a match, and the root is the final. Seed placement follows the doubling recurrence

$$S(1) = [1], \quad S(2k) = [s, 2k + 1 - s : s \in S(k)],$$

giving $S(8) = [1, 8, 4, 5, 2, 7, 3, 6]$. Every first round pair sums to $k + 1$ (the strongest seed against the weakest), and the top 2^t seeds land in 2^t different subtrees. Figure 2 shows the eight seed tree with the paths of seeds 1 and 2 highlighted.

The placement's fairness guarantee can be stated and proved exactly.

Lemma 1 (Subtree isolation). *In the bracket built from $S(k)$, every subtree with 2^j leaves contains exactly one of the top $k/2^j$ seeds.*




Figure 2: The seeded eight player tree, seed order $[1, 8, 4, 5, 2, 7, 3, 6]$. The paths of seeds 1 and 2 (ringed) stay in opposite halves and meet only at the root, so seeds 1 and 2 can meet only in the final.

Proof. Induct on the doubling. For $S(2) = [1, 2]$ the claim is immediate. Assume it holds for $S(m)$ and build $S(2m)$. The size two subtrees of the new bracket are exactly the pairs $(s, 2m+1-s)$ with $s \in S(m)$, and since $s \leq m < 2m+1-s$, each contains exactly one of the top m seeds; that is the $j = 1$ case. Now contract every pair to its smaller member. By construction this turns the $S(2m)$ bracket into precisely the $S(m)$ bracket, and a subtree with 2^j leaves ($j \geq 1$) contracts to one with 2^{j-1} leaves, which by the induction hypothesis contains exactly one of the top $m/2^{j-1} = 2m/2^j$ seeds. The discarded partners all exceed m , hence also $2m/2^j$, so the count is unchanged. $\square$

Two players meet exactly when the knockout reaches the root of the smallest subtree containing both, so the lemma pins down who can meet when: the top 2^t seeds occupy 2^t different subtrees of size $k/2^t$, hence no two of them can meet before the round with 2^t players left. In particular seeds 1 and 2 can meet only in the final, and seeds 1 through 4 no earlier than the semifinals. The application computes the exact earliest meeting round as `seedMeetRound`, the height of the two leaves' lowest common ancestor in the tree. A single elimination of k players is exactly $k - 1$ matches, since each match produces one loss and $k - 1$ players must be eliminated once.

4.2 Double elimination and the second life

Single elimination is short and gives a clean tree, but one unlucky match, a bad roll or a lost break, ends a strong player's run. Double elimination fixes this and suits a casual event, because everyone plays more. Each player starts in the winners bracket; a first loss drops them into the losers bracket, and a second loss ends their run. So every player except the champion is eliminated with exactly two losses. Figure 3 shows the flow for four players.

Counting by losses gives the totals directly. The $k - 1$ non champions each carry two losses, that is $2(k - 1)$ losses, and the champion carries 0 or 1, so the number of matches is

$$2(k - 1) + \{0 \text{ or } 1\} = \begin{cases} 2k - 2 & \text{no reset} \\ 2k - 1 & \text{grand final reset.} \end{cases}$$

The reset arises because the WB champion enters the grand final undefeated: if the LB champion wins the first grand final game, both then have one loss, so a single deciding game is played. For $k = 8$ this is 14 or 15 matches against 7 for single elimination. Table 2 summarizes the tradeoff.




Figure 3: Double elimination for $k = 4$. A winners bracket (WB) loss (edge L) drops a player into the losers bracket (LB); a second loss ends the run. The undefeated WB champion meets the once beaten LB champion in the grand final.

	single elimination	double elimination
eliminated after	1 loss	2 losses
matches for k seeds	$k - 1$	$2k - 2$ (reset: $2k - 1$)
rounds for $k = 8$	3	8
minimum matches per player	1	2
survives one unlucky match	no	yes

Table 2: The knockout tradeoff. Double elimination roughly doubles the play and guarantees every entrant at least two matches, at the cost of a longer event; the application offers both, chosen with a format toggle.

The construction `buildDoubleBracket` handles any $k = 2^m$. The winners bracket is the ordinary seeded tree ($k - 1$ matches). The losers bracket has $2(m - 1)$ rounds ($k - 2$ matches) that alternate minor rounds, where losers bracket survivors meet, and major rounds, where fresh winners bracket losers drop in. The drops are crossed in reversed order, so a dropping player meets a survivor from the opposite side of the draw rather than someone from their own quarter. This avoids rematches for as long as the bracket is wide enough to permit it; once the losers bracket narrows to single matches, rematches can occur, as in any double elimination (for $k = 4$ they are already unavoidable, since the losers final has only one possible pairing). The whole bracket is built as a list of matches already interleaved into a valid play order, each carrying a pointer to where its winner and its loser go next, so a single match queue drives the simulation. For $k = 8$ ($m = 3$) the losers bracket rounds hold 2, 2, 1, 1 matches, six in all ($k - 2$), so a full run is 7 winners bracket matches, 6 losers bracket matches, and one grand final, that is $14 = 2k - 2$.

5 A documented limitation: the carry-over effect

If player i faces j in one round and k in the next, then k is said to receive a *carry-over* from j through i , for example inheriting an opponent who is tired or on a run. Reading each player's opponent sequence and wrapping the last round to the first, one tallies the matrix G where $G[j][k]$ counts these handoffs, and measures imbalance by the sum of squares $\Gamma = \sum_{j,k} G[j][k]^2$. A balanced schedule spreads the handoffs so every off diagonal entry is 1, reaching the minimum $\Gamma = n(n - 1)$.

The circle method does the opposite of balancing. Russell (1980) showed a balanced schedule exists whenever n is a power of two, using finite field constructions, while Lambrechts, Ficker, Goossens,




Figure 4: The carry-over matrix of the circle method schedule of K_6 . The fixed player (row 0) hands off evenly, but the ring players concentrate into cells of weight 3. Its value is $\Gamma = 60$, exactly twice the balanced minimum $n(n-1) = 30$.

and Spieksma (2018) proved that the circle method attains the *maximum* possible carry-over value, and that every schedule achieving that maximum is a circle method schedule. So the very rotation that made the round robin easy to build and prove is, provably, the least balanced choice for carry-over. For a casual Student Union event the trade for simplicity is acceptable, and the application states the limitation and draws the matrix live (Figure 4). This honesty, showing a weakness of the chosen method with a citation, is itself part of the design.

The extremes are visible on small fields. For $n = 4$ the complete graph K_4 has a unique one-factorization, so the circle method has no freedom and its schedule is already balanced: every off diagonal entry of G is 1 and $\Gamma = 12 = n(n-1)$, the minimum. The gap between the circle method’s maximum and the balanced minimum opens only for $n \geq 6$, where the circle schedule of K_6 scores $\Gamma = 60$, exactly twice the balanced minimum of 30 (Figure 4). Russell’s balanced designs, available whenever n is a power of two, recover the minimum at the cost of the simple rotation and its one line proof, which is the tradeoff a real event has to weigh.

6 Presenting the mathematics: the guided walkthrough

Because the point of the project is the mathematics, the application does not just use these results, it teaches them. A built in *guided walkthrough* presents the two core algorithms in eighteen steps across four acts (round robin, rating, knockout tree, carry-over), each step pairing a short narration with a live visual: the K_6 diagram inking one matching per rotation with chalk residue, the congruence table of Table 1 computed from the shipped scheduler (Figure 5), the doubling seed order built pass by pass, the meeting round highlighted on the tree, and the carry-over heatmap. The walkthrough runs on a fixed teaching field, independent of the loaded roster, so it presents identically every time, and it advances by button or arrow key, which is how the ten minute live demonstration of this project is delivered: the walkthrough carries the theory, and the live tournament run carries the practice.

Guided Walkthrough

The two discrete-math algorithms behind the tournament, one step at a time, worked on a small field so every number stays on screen. Click Next, or drive it with the left and right arrow keys.

ROUND ROBIN RATING KNOCKOUT TREE CARRY-OVER Step 9 of 18

ROUND ROBIN, THE PROOF

Why every pair meets exactly once

Rotating felt right, but why is it guaranteed? Number the five ring players 1 to 5. Opposite seats always have a constant position sum, and each rotation shifts that sum by a fixed step, so:

ring players i and j meet in round $r+1$ exactly when $i + j = -2r \pmod{5}$. Read the residue column below: the five rounds hit the classes 0, 3, 1, 4, 2, which is every class once.

Why once and not twice? Because 5 is odd, **2 is invertible mod 5** ($2 \times 3 = 6 \equiv 1$), so $r \mapsto -2r$ is a bijection. Each ring pair lands in exactly one round; player 0 takes the leftover.

5 odd $\Rightarrow$ 2 invertible in $\mathbb{Z}_5$, since $2 \times 3 \equiv 1$
 $r \mapsto -2r$ is a bijection $\Rightarrow$ **every pair, exactly one round**

R	$-2R \pmod{5}$	RING PAIR, SUM CONGRUENT	FIXED
0	0	{1,4} {2,3}	0 vs 5
1	3	{5,3} {1,2}	0 vs 4
2	1	{4,2} {5,1}	0 vs 3
3	4	{3,1} {4,5}	0 vs 2
4	2	{2,5} {3,4}	0 vs 1

BACK AUTO-PLAY RESTART NEXT

Figure 5: The walkthrough at the proof step of Act 1. The congruence table on the right is generated live from the shipped `buildSchedule`, so the displayed proof and the running code cannot disagree.

7 Verification and results

Correctness is established two ways. First, the application’s algorithms live in a DOM free block that a Node test runner extracts and exercises directly, so the tested code is the shipped code. The suite runs roughly 108,000 assertions: the circle method’s structure for every field size $n = 4$ to 30, the congruence invariant of Theorem 1, the Elo identities, the race inversion round trip and its Monte Carlo, the seed order and meeting rounds, the carry-over totals, and full simulated double eliminations that confirm the two loss invariant and the $2k - 2$ or $2k - 1$ match count for $k = 4, 8, 16$. Second, the in application verification panel recomputes the round robin’s five structural properties from the generated schedule on every run.

Figure 6 shows a completed double elimination the application played for a field of thirteen (top eight seeded). The run displays every piece of the mathematics at once: the eighth seed upsets the first in the opening round, first losses drop into the losers bracket and climb back through its alternating rounds, the losers champion takes the first grand final game, and the reset game decides the title, $15 = 2k - 1$ matches in all. Fields up to thirty players keep every guarantee: at $n = 30$ the group stage is 29 rounds and 435 matches, all five checks pass, and Elo is conserved.

8 Design decisions

- **Graph and one-factorization framing.** It turns vague scheduling into provable structure, which is the point of the course.
- **Circle method for the round robin.** The simplest construction with a short, fully rigorous proof. The honest cost is maximum carry-over, which the report and the app both state.
- **Calibrated race model over raw Elo as rack probability.** Keeps match odds faithful to ratings while producing realistic scorelines.



Figure 6: A completed double elimination in the application: the winners bracket (the seeded tree), the losers bracket (six matches over four rounds for $k = 8$), and a grand final that went to the reset game.

- **Both knockout formats.** Single elimination gives the clean binary tree for the mathematics; double elimination is fairer to one unlucky match and keeps everyone playing, which fits the event. Offering both makes the $k - 1$ versus $2k - 2$ tradeoff visible.
- **Live verification and an extracted test block.** Correctness is demonstrated from output, not asserted.

9 What I learned

The main lesson was how much of an apparently practical problem is really a question about discrete structures, and how much that reframing pays off. Once who plays whom became the edges of K_n , the schedule's correctness reduced to a one line fact about modular arithmetic, and its fairness properties, byes and carry-over, became theorems rather than guesses. I came to value constructions that arrive with proofs: the circle method was chosen not because it is the only schedule (there are 6240 for eight players) but because its correctness is short to state and to defend, which matters when a professor can ask about it directly.

I also learned the discipline of showing a method's weakness. The carry-over result, that my own chosen schedule is provably the worst for balance, was more convincing to include than to hide, and stating it forced me to actually understand the tradeoff rather than assume the method was ideal. Building the double elimination taught me to think in invariants: instead of tracking a tangle of

two brackets, the entire format is pinned down by one counting argument, every non champion leaves with exactly two losses, which both guided the implementation and became the property the tests check. Finally, keeping the simulation honest to the ratings, by inverting the race formula rather than reusing the Elo number, was a small reminder that a model can be subtly wrong even when each piece looks right.

10 Conclusion

Behind a pool tournament sits a compact and provable core of discrete mathematics: a one-factorization of K_n built by a rotation that works because 2 is a unit modulo $n - 1$, a chromatic index that explains byes, a negative binomial race calibrated to Elo, seeded binary trees with an exact meeting distance, a loss counting argument that pins down double elimination at $2k - 2$ matches, and a carry-over analysis in which the chosen method is provably extremal. The system runs the tournament and, at every step, shows and checks the mathematics that makes its decisions correct and fair.

References

- [1] T. P. Kirkman, On a problem in combinations, *Cambridge and Dublin Mathematical Journal* 2 (1847) 191-204.
- [2] É. Lucas, *Récréations Mathématiques*, vol. 2, Gauthier-Villars, Paris, 1883.
- [3] K. G. Russell, Balancing carry-over effects in round robin tournaments, *Biometrika* 67 (1980) 127-131.
- [4] E. Lambrechts, A. M. C. Ficker, D. R. Goossens, F. C. R. Spieksma, Round-robin tournaments generated by the circle method have maximum carry-over, *Mathematical Programming* 172 (2018) 277-302.
- [5] A. E. Elo, *The Rating of Chessplayers, Past and Present*, Arco, 1978.
- [6] W. D. Wallis, *One-Factorizations*, Kluwer, 1997.
- [7] V. G. Vizing, On an estimate of the chromatic class of a p-graph, *Diskret. Analiz* 3 (1964) 25-30.
- [8] OEIS Foundation, Sequence A000438, Number of one-factorizations of K_{2n} .